

# Overfitting and Underfitting

June 16, 2019

## 1 Overfitting and underfitting

**Overfit** When the model adapts to the particularities of the training set. We want to generalize it for better performance in unknown data.

It can happen when: - we train the model for too long (the validation accuracy hits a peak and starts decreasing) - the model has too much freedom (too many layers or hidden units)

Solutions: 1. Train on large datasets 2. Use regularization 3. Find the best training duration

**Underfit** When there is still room for improvement on test data.

It can happen when: - the model is not powerful enough - it is overregularized - was not trained enough

```
[1]: from __future__ import absolute_import, division, print_function

import tensorflow as tf
from tensorflow import keras

import numpy as np
import matplotlib.pyplot as plt

print(tf.__version__)
```

1.13.1

### 1.1 Preprocessing

We will use the IMDB review dataset with multi-hot encoding ([3, 5] => [000101000...00]) instead of embedding. The model will quickly overfit.

```
[2]: NUM_WORDS = 10000

(train_data, train_labels), (test_data, test_labels) = keras.datasets.imdb.
    ↳load_data(num_words=NUM_WORDS)

def multi_hot_sequences(sequences, dimension):
    # Create an all-zero matrix of shape (len(sequences), dimension)
    results = np.zeros((len(sequences), dimension))
    for i, word_indices in enumerate(sequences):
```

```

        results[i, word_indices] = 1.0 # set specific indices of results[i] to 1
→1s
    return results

```

```


train_data = multi_hot_sequences(train_data, dimension=NUM_WORDS)
test_data = multi_hot_sequences(test_data, dimension=NUM_WORDS)

```

Since the words are sorted by frequency, we expect more 1's near index 0

```
[3]: plt.plot(train_data[0])
```

```
[3]: [<matplotlib.lines.Line2D at 0x2d2985021d0>]
```



## 1.2 About overfitting

The simplest way to prevent overfitting is to reduce the size (**capacity** = number of learnable parameters = number of layers and units) so that it will focus on more important patterns with more predictive power.

To find the best architecture, the best is to start with just a few layers and units and then increase them until validation loss stops improving.

### Baseline model

```

[4]: baseline_model = keras.Sequential([
    # `input_shape` is only required here so that `.summary` works.
    keras.layers.Dense(16, activation=tf.nn.relu, input_shape=(NUM_WORDS,)),
    keras.layers.Dense(16, activation=tf.nn.relu),

```

```

        keras.layers.Dense(1, activation=tf.nn.sigmoid)
    ])

baseline_model.compile(optimizer='adam',
                      loss='binary_crossentropy',
                      metrics=['accuracy', 'binary_crossentropy'])

baseline_model.summary()

```

WARNING:tensorflow:From C:\Users\Pedro\AppData\Roaming\Python\Python37\site-packages\tensorflow\python\ops\resource\_variable\_ops.py:435: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:  
Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	160016
dense_1 (Dense)	(None, 16)	272
dense_2 (Dense)	(None, 1)	17

Total params: 160,305  
 Trainable params: 160,305  
 Non-trainable params: 0

```

[5]: baseline_history = baseline_model.fit(train_data,
                                           train_labels,
                                           epochs=20,
                                           batch_size=512,
                                           validation_data=(test_data, test_labels),
                                           verbose=2)

```

Train on 25000 samples, validate on 25000 samples

WARNING:tensorflow:From C:\Users\Pedro\AppData\Roaming\Python\Python37\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:  
Use tf.cast instead.

Epoch 1/20

- 19s - loss: 0.4814 - acc: 0.8046 - binary\_crossentropy: 0.4814 - val\_loss: 0.3301 - val\_acc: 0.8766 - val\_binary\_crossentropy: 0.3301

Epoch 2/20

- 15s - loss: 0.2442 - acc: 0.9126 - binary\_crossentropy: 0.2442 - val\_loss: 0.2835 - val\_acc: 0.8885 - val\_binary\_crossentropy: 0.2835  
Epoch 3/20  
- 11s - loss: 0.1777 - acc: 0.9376 - binary\_crossentropy: 0.1777 - val\_loss: 0.2952 - val\_acc: 0.8828 - val\_binary\_crossentropy: 0.2952  
Epoch 4/20  
- 8s - loss: 0.1449 - acc: 0.9506 - binary\_crossentropy: 0.1449 - val\_loss: 0.3374 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.3374  
Epoch 5/20  
- 8s - loss: 0.1193 - acc: 0.9604 - binary\_crossentropy: 0.1193 - val\_loss: 0.3429 - val\_acc: 0.8736 - val\_binary\_crossentropy: 0.3429  
Epoch 6/20  
- 7s - loss: 0.0981 - acc: 0.9695 - binary\_crossentropy: 0.0981 - val\_loss: 0.3738 - val\_acc: 0.8693 - val\_binary\_crossentropy: 0.3738  
Epoch 7/20  
- 7s - loss: 0.0817 - acc: 0.9753 - binary\_crossentropy: 0.0817 - val\_loss: 0.4114 - val\_acc: 0.8656 - val\_binary\_crossentropy: 0.4114  
Epoch 8/20  
- 7s - loss: 0.0699 - acc: 0.9806 - binary\_crossentropy: 0.0699 - val\_loss: 0.4512 - val\_acc: 0.8621 - val\_binary\_crossentropy: 0.4512  
Epoch 9/20  
- 7s - loss: 0.0563 - acc: 0.9858 - binary\_crossentropy: 0.0563 - val\_loss: 0.4873 - val\_acc: 0.8606 - val\_binary\_crossentropy: 0.4873  
Epoch 10/20  
- 7s - loss: 0.0475 - acc: 0.9889 - binary\_crossentropy: 0.0475 - val\_loss: 0.5306 - val\_acc: 0.8577 - val\_binary\_crossentropy: 0.5306  
Epoch 11/20  
- 12s - loss: 0.0383 - acc: 0.9928 - binary\_crossentropy: 0.0383 - val\_loss: 0.5761 - val\_acc: 0.8559 - val\_binary\_crossentropy: 0.5761  
Epoch 12/20  
- 10s - loss: 0.0308 - acc: 0.9946 - binary\_crossentropy: 0.0308 - val\_loss: 0.6153 - val\_acc: 0.8537 - val\_binary\_crossentropy: 0.6153  
Epoch 13/20  
- 7s - loss: 0.0255 - acc: 0.9963 - binary\_crossentropy: 0.0255 - val\_loss: 0.6728 - val\_acc: 0.8514 - val\_binary\_crossentropy: 0.6728  
Epoch 14/20  
- 6s - loss: 0.0199 - acc: 0.9978 - binary\_crossentropy: 0.0199 - val\_loss: 0.7107 - val\_acc: 0.8506 - val\_binary\_crossentropy: 0.7107  
Epoch 15/20  
- 6s - loss: 0.0157 - acc: 0.9984 - binary\_crossentropy: 0.0157 - val\_loss: 0.7351 - val\_acc: 0.8498 - val\_binary\_crossentropy: 0.7351  
Epoch 16/20  
- 7s - loss: 0.0125 - acc: 0.9991 - binary\_crossentropy: 0.0125 - val\_loss: 0.7775 - val\_acc: 0.8498 - val\_binary\_crossentropy: 0.7775  
Epoch 17/20  
- 7s - loss: 0.0103 - acc: 0.9994 - binary\_crossentropy: 0.0103 - val\_loss: 0.8070 - val\_acc: 0.8488 - val\_binary\_crossentropy: 0.8070  
Epoch 18/20

```

- 6s - loss: 0.0084 - acc: 0.9996 - binary_crossentropy: 0.0084 - val_loss:
0.8404 - val_acc: 0.8485 - val_binary_crossentropy: 0.8404
Epoch 19/20
- 6s - loss: 0.0069 - acc: 0.9996 - binary_crossentropy: 0.0069 - val_loss:
0.8704 - val_acc: 0.8485 - val_binary_crossentropy: 0.8704
Epoch 20/20
- 6s - loss: 0.0057 - acc: 0.9997 - binary_crossentropy: 0.0057 - val_loss:
0.8976 - val_acc: 0.8476 - val_binary_crossentropy: 0.8976

```

### Smaller model

```

[6]: smaller_model = keras.Sequential([
    keras.layers.Dense(4, activation=tf.nn.relu, input_shape=(NUM_WORDS,)),
    keras.layers.Dense(4, activation=tf.nn.relu),
    keras.layers.Dense(1, activation=tf.nn.sigmoid)
])

smaller_model.compile(optimizer='adam',
                      loss='binary_crossentropy',
                      metrics=['accuracy', 'binary_crossentropy'])

smaller_model.summary()

```

```

-----
Layer (type)                 Output Shape              Param #
-----
dense_3 (Dense)              (None, 4)                 40004
-----
dense_4 (Dense)              (None, 4)                 20
-----
dense_5 (Dense)              (None, 1)                 5
=====
Total params: 40,029
Trainable params: 40,029
Non-trainable params: 0
-----

```

```

[7]: smaller_history = smaller_model.fit(train_data,
                                         train_labels,
                                         epochs=20,
                                         batch_size=512,
                                         validation_data=(test_data, test_labels),
                                         verbose=2)

```

Train on 25000 samples, validate on 25000 samples

```

Epoch 1/20
- 6s - loss: 0.6613 - acc: 0.5813 - binary_crossentropy: 0.6613 - val_loss:
0.6218 - val_acc: 0.6428 - val_binary_crossentropy: 0.6218

```

Epoch 2/20  
- 6s - loss: 0.5774 - acc: 0.7418 - binary\_crossentropy: 0.5774 - val\_loss: 0.5550 - val\_acc: 0.7581 - val\_binary\_crossentropy: 0.5550  
Epoch 3/20  
- 6s - loss: 0.5126 - acc: 0.8280 - binary\_crossentropy: 0.5126 - val\_loss: 0.5091 - val\_acc: 0.8053 - val\_binary\_crossentropy: 0.5091  
Epoch 4/20  
- 6s - loss: 0.4657 - acc: 0.8676 - binary\_crossentropy: 0.4657 - val\_loss: 0.4767 - val\_acc: 0.8368 - val\_binary\_crossentropy: 0.4767  
Epoch 5/20  
- 6s - loss: 0.4300 - acc: 0.8933 - binary\_crossentropy: 0.4300 - val\_loss: 0.4560 - val\_acc: 0.8439 - val\_binary\_crossentropy: 0.4560  
Epoch 6/20  
- 6s - loss: 0.4015 - acc: 0.9086 - binary\_crossentropy: 0.4015 - val\_loss: 0.4361 - val\_acc: 0.8650 - val\_binary\_crossentropy: 0.4361  
Epoch 7/20  
- 6s - loss: 0.3772 - acc: 0.9203 - binary\_crossentropy: 0.3772 - val\_loss: 0.4238 - val\_acc: 0.8692 - val\_binary\_crossentropy: 0.4238  
Epoch 8/20  
- 6s - loss: 0.3563 - acc: 0.9313 - binary\_crossentropy: 0.3563 - val\_loss: 0.4176 - val\_acc: 0.8661 - val\_binary\_crossentropy: 0.4176  
Epoch 9/20  
- 5s - loss: 0.3376 - acc: 0.9391 - binary\_crossentropy: 0.3376 - val\_loss: 0.4125 - val\_acc: 0.8662 - val\_binary\_crossentropy: 0.4125  
Epoch 10/20  
- 5s - loss: 0.3200 - acc: 0.9460 - binary\_crossentropy: 0.3200 - val\_loss: 0.4012 - val\_acc: 0.8716 - val\_binary\_crossentropy: 0.4012  
Epoch 11/20  
- 6s - loss: 0.2960 - acc: 0.9521 - binary\_crossentropy: 0.2960 - val\_loss: 0.3901 - val\_acc: 0.8725 - val\_binary\_crossentropy: 0.3901  
Epoch 12/20  
- 6s - loss: 0.2582 - acc: 0.9584 - binary\_crossentropy: 0.2582 - val\_loss: 0.3659 - val\_acc: 0.8762 - val\_binary\_crossentropy: 0.3659  
Epoch 13/20  
- 6s - loss: 0.2134 - acc: 0.9636 - binary\_crossentropy: 0.2134 - val\_loss: 0.3525 - val\_acc: 0.8750 - val\_binary\_crossentropy: 0.3525  
Epoch 14/20  
- 6s - loss: 0.1540 - acc: 0.9683 - binary\_crossentropy: 0.1540 - val\_loss: 0.3330 - val\_acc: 0.8749 - val\_binary\_crossentropy: 0.3330  
Epoch 15/20  
- 6s - loss: 0.1172 - acc: 0.9713 - binary\_crossentropy: 0.1172 - val\_loss: 0.3657 - val\_acc: 0.8711 - val\_binary\_crossentropy: 0.3657  
Epoch 16/20  
- 6s - loss: 0.1006 - acc: 0.9757 - binary\_crossentropy: 0.1006 - val\_loss: 0.3760 - val\_acc: 0.8706 - val\_binary\_crossentropy: 0.3760  
Epoch 17/20  
- 6s - loss: 0.0882 - acc: 0.9801 - binary\_crossentropy: 0.0882 - val\_loss: 0.3924 - val\_acc: 0.8690 - val\_binary\_crossentropy: 0.3924

```
Epoch 18/20
- 6s - loss: 0.0788 - acc: 0.9825 - binary_crossentropy: 0.0788 - val_loss:
0.4186 - val_acc: 0.8691 - val_binary_crossentropy: 0.4186
Epoch 19/20
- 6s - loss: 0.0709 - acc: 0.9850 - binary_crossentropy: 0.0709 - val_loss:
0.4316 - val_acc: 0.8670 - val_binary_crossentropy: 0.4316
Epoch 20/20
- 6s - loss: 0.0640 - acc: 0.9870 - binary_crossentropy: 0.0640 - val_loss:
0.4495 - val_acc: 0.8664 - val_binary_crossentropy: 0.4495
```

### Bigger model

```
[8]: bigger_model = keras.models.Sequential([
    keras.layers.Dense(512, activation=tf.nn.relu, input_shape=(NUM_WORDS,)),
    keras.layers.Dense(512, activation=tf.nn.relu),
    keras.layers.Dense(1, activation=tf.nn.sigmoid)
])

bigger_model.compile(optimizer='adam',
                    loss='binary_crossentropy',
                    metrics=['accuracy', 'binary_crossentropy'])

bigger_model.summary()
```

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 512)	5120512
dense_7 (Dense)	(None, 512)	262656
dense_8 (Dense)	(None, 1)	513

Total params: 5,383,681  
 Trainable params: 5,383,681  
 Non-trainable params: 0

```
[9]: bigger_history = bigger_model.fit(train_data, train_labels,
                                       epochs=20,
                                       batch_size=512,
                                       validation_data=(test_data, test_labels),
                                       verbose=2)
```

Train on 25000 samples, validate on 25000 samples

```
Epoch 1/20
- 18s - loss: 0.3461 - acc: 0.8520 - binary_crossentropy: 0.3461 - val_loss:
0.2922 - val_acc: 0.8817 - val_binary_crossentropy: 0.2922
```

Epoch 2/20  
 - 14s - loss: 0.1437 - acc: 0.9478 - binary\_crossentropy: 0.1437 - val\_loss: 0.3330 - val\_acc: 0.8723 - val\_binary\_crossentropy: 0.3330

Epoch 3/20  
 - 13s - loss: 0.0465 - acc: 0.9872 - binary\_crossentropy: 0.0465 - val\_loss: 0.4505 - val\_acc: 0.8696 - val\_binary\_crossentropy: 0.4505

Epoch 4/20  
 - 13s - loss: 0.0072 - acc: 0.9986 - binary\_crossentropy: 0.0072 - val\_loss: 0.5805 - val\_acc: 0.8685 - val\_binary\_crossentropy: 0.5805

Epoch 5/20  
 - 14s - loss: 9.1697e-04 - acc: 1.0000 - binary\_crossentropy: 9.1697e-04 - val\_loss: 0.6804 - val\_acc: 0.8703 - val\_binary\_crossentropy: 0.6804

Epoch 6/20  
 - 16s - loss: 2.4898e-04 - acc: 1.0000 - binary\_crossentropy: 2.4898e-04 - val\_loss: 0.7200 - val\_acc: 0.8706 - val\_binary\_crossentropy: 0.7200

Epoch 7/20  
 - 13s - loss: 1.5187e-04 - acc: 1.0000 - binary\_crossentropy: 1.5187e-04 - val\_loss: 0.7462 - val\_acc: 0.8705 - val\_binary\_crossentropy: 0.7462

Epoch 8/20  
 - 14s - loss: 1.0835e-04 - acc: 1.0000 - binary\_crossentropy: 1.0835e-04 - val\_loss: 0.7670 - val\_acc: 0.8704 - val\_binary\_crossentropy: 0.7670

Epoch 9/20  
 - 13s - loss: 8.1950e-05 - acc: 1.0000 - binary\_crossentropy: 8.1950e-05 - val\_loss: 0.7837 - val\_acc: 0.8706 - val\_binary\_crossentropy: 0.7837

Epoch 10/20  
 - 13s - loss: 6.3951e-05 - acc: 1.0000 - binary\_crossentropy: 6.3951e-05 - val\_loss: 0.7994 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.7994

Epoch 11/20  
 - 14s - loss: 5.0951e-05 - acc: 1.0000 - binary\_crossentropy: 5.0951e-05 - val\_loss: 0.8140 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8140

Epoch 12/20  
 - 13s - loss: 4.1561e-05 - acc: 1.0000 - binary\_crossentropy: 4.1561e-05 - val\_loss: 0.8266 - val\_acc: 0.8707 - val\_binary\_crossentropy: 0.8266

Epoch 13/20  
 - 13s - loss: 3.4086e-05 - acc: 1.0000 - binary\_crossentropy: 3.4086e-05 - val\_loss: 0.8384 - val\_acc: 0.8709 - val\_binary\_crossentropy: 0.8384

Epoch 14/20  
 - 13s - loss: 2.8378e-05 - acc: 1.0000 - binary\_crossentropy: 2.8378e-05 - val\_loss: 0.8502 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8502

Epoch 15/20  
 - 13s - loss: 2.3815e-05 - acc: 1.0000 - binary\_crossentropy: 2.3815e-05 - val\_loss: 0.8611 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8611

Epoch 16/20  
 - 14s - loss: 2.0211e-05 - acc: 1.0000 - binary\_crossentropy: 2.0211e-05 - val\_loss: 0.8716 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8716

Epoch 17/20  
 - 13s - loss: 1.7215e-05 - acc: 1.0000 - binary\_crossentropy: 1.7215e-05 - val\_loss: 0.8816 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8816



Epoch 18/20

- 13s - loss: 1.4778e-05 - acc: 1.0000 - binary\_crossentropy: 1.4778e-05 -  
val\_loss: 0.8910 - val\_acc: 0.8708 - val\_binary\_crossentropy: 0.8910

Epoch 19/20

- 14s - loss: 1.2767e-05 - acc: 1.0000 - binary\_crossentropy: 1.2767e-05 -  
val\_loss: 0.9000 - val\_acc: 0.8706 - val\_binary\_crossentropy: 0.9000

Epoch 20/20

- 13s - loss: 1.1097e-05 - acc: 1.0000 - binary\_crossentropy: 1.1097e-05 -  
val\_loss: 0.9085 - val\_acc: 0.8706 - val\_binary\_crossentropy: 0.9085

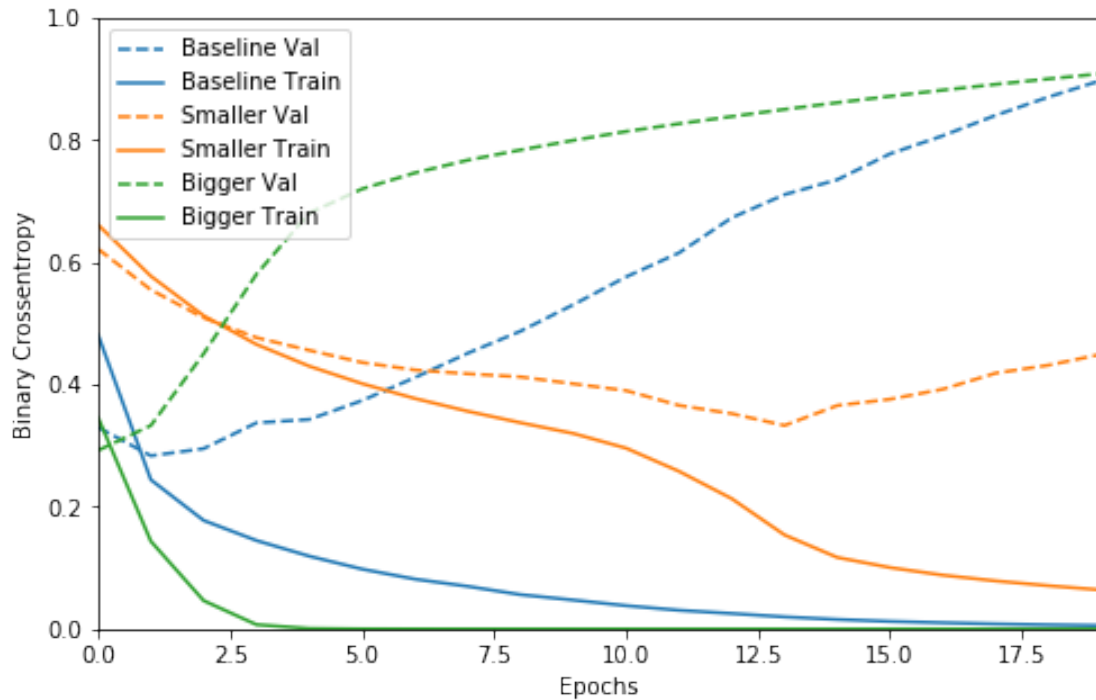
```
[17]: def plot_history(histories, key='binary_crossentropy', y_min=0):
    plt.figure(figsize=(8,5))

    for name, history in histories:
        val = plt.plot(history.epoch, history.history['val_'+key],
                        '--', label=name.title()+' Val')
        plt.plot(history.epoch, history.history[key], color=val[0].get_color(),
                 label=name.title()+' Train')

    plt.xlabel('Epochs')
    plt.ylabel(key.replace('_', ' ').title())
    plt.legend()

    plt.xlim([0,max(history.epoch)])
    plt.ylim([y_min, 1])

plot_history([('baseline', baseline_history),
             ('smaller', smaller_history),
             ('bigger', bigger_history)])
```



The bigger network starts to overfit after one epoch and much more severely. The smaller starts to overfit after the baseline and more slowly. So more capacity, means low training loss and high susceptibility to overfit (large difference between training and validation loss).

## 1.3 Strategies

### 1.3.1 Weight regularization

The simpler models are usually the best against overfitting. So if it has less parameters or if they have less entropy it is simpler. Thus forcing the weights to take small values, makes the distribution more “regular”.

This is done by adding to the loss function a cost associated to large weights. There are two possibilities: - L1 regularization: cost added is proportional to the weights (L1 norms) - L2 regularization (weight decay): cost added is proportional to the square of the weights (L2 norms)

```
[11]: l2_model = keras.models.Sequential([
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),
        activation=tf.nn.relu, input_shape=(NUM_WORDS,)),
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),
        activation=tf.nn.relu),
    keras.layers.Dense(1, activation=tf.nn.sigmoid)
])

l2_model.compile(optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy', 'binary_crossentropy'])
```

```
l2_model_history = l2_model.fit(train_data, train_labels,
                                epochs=20,
                                batch_size=512,
                                validation_data=(test_data, test_labels),
                                verbose=2)
```

Train on 25000 samples, validate on 25000 samples

Epoch 1/20

- 6s - loss: 0.5198 - acc: 0.8112 - binary\_crossentropy: 0.4784 - val\_loss: 0.3732 - val\_acc: 0.8774 - val\_binary\_crossentropy: 0.3301

Epoch 2/20

- 5s - loss: 0.2977 - acc: 0.9102 - binary\_crossentropy: 0.2501 - val\_loss: 0.3349 - val\_acc: 0.8878 - val\_binary\_crossentropy: 0.2841

Epoch 3/20

- 5s - loss: 0.2475 - acc: 0.9320 - binary\_crossentropy: 0.1942 - val\_loss: 0.3401 - val\_acc: 0.8849 - val\_binary\_crossentropy: 0.2850

Epoch 4/20

- 5s - loss: 0.2269 - acc: 0.9418 - binary\_crossentropy: 0.1699 - val\_loss: 0.3599 - val\_acc: 0.8784 - val\_binary\_crossentropy: 0.3015

Epoch 5/20

- 5s - loss: 0.2117 - acc: 0.9492 - binary\_crossentropy: 0.1522 - val\_loss: 0.3705 - val\_acc: 0.8762 - val\_binary\_crossentropy: 0.3100

Epoch 6/20

- 5s - loss: 0.2000 - acc: 0.9544 - binary\_crossentropy: 0.1384 - val\_loss: 0.3836 - val\_acc: 0.8749 - val\_binary\_crossentropy: 0.3213

Epoch 7/20

- 5s - loss: 0.1924 - acc: 0.9577 - binary\_crossentropy: 0.1291 - val\_loss: 0.4065 - val\_acc: 0.8711 - val\_binary\_crossentropy: 0.3423

Epoch 8/20

- 5s - loss: 0.1856 - acc: 0.9596 - binary\_crossentropy: 0.1205 - val\_loss: 0.4195 - val\_acc: 0.8698 - val\_binary\_crossentropy: 0.3536

Epoch 9/20

- 5s - loss: 0.1790 - acc: 0.9624 - binary\_crossentropy: 0.1120 - val\_loss: 0.4389 - val\_acc: 0.8647 - val\_binary\_crossentropy: 0.3711

Epoch 10/20

- 5s - loss: 0.1719 - acc: 0.9665 - binary\_crossentropy: 0.1035 - val\_loss: 0.4474 - val\_acc: 0.8663 - val\_binary\_crossentropy: 0.3783

Epoch 11/20

- 5s - loss: 0.1641 - acc: 0.9704 - binary\_crossentropy: 0.0946 - val\_loss: 0.4612 - val\_acc: 0.8653 - val\_binary\_crossentropy: 0.3910

Epoch 12/20

- 5s - loss: 0.1585 - acc: 0.9734 - binary\_crossentropy: 0.0878 - val\_loss: 0.4777 - val\_acc: 0.8626 - val\_binary\_crossentropy: 0.4066

Epoch 13/20

- 5s - loss: 0.1517 - acc: 0.9768 - binary\_crossentropy: 0.0802 - val\_loss: 0.4913 - val\_acc: 0.8614 - val\_binary\_crossentropy: 0.4195


```

Epoch 14/20
- 5s - loss: 0.1486 - acc: 0.9773 - binary_crossentropy: 0.0762 - val_loss:
0.5088 - val_acc: 0.8591 - val_binary_crossentropy: 0.4360
Epoch 15/20
- 5s - loss: 0.1444 - acc: 0.9799 - binary_crossentropy: 0.0714 - val_loss:
0.5186 - val_acc: 0.8588 - val_binary_crossentropy: 0.4452
Epoch 16/20
- 5s - loss: 0.1409 - acc: 0.9808 - binary_crossentropy: 0.0668 - val_loss:
0.5322 - val_acc: 0.8578 - val_binary_crossentropy: 0.4579
Epoch 17/20
- 5s - loss: 0.1382 - acc: 0.9819 - binary_crossentropy: 0.0641 - val_loss:
0.5494 - val_acc: 0.8566 - val_binary_crossentropy: 0.4749
Epoch 18/20
- 5s - loss: 0.1371 - acc: 0.9826 - binary_crossentropy: 0.0619 - val_loss:
0.5586 - val_acc: 0.8564 - val_binary_crossentropy: 0.4829
Epoch 19/20
- 5s - loss: 0.1317 - acc: 0.9842 - binary_crossentropy: 0.0563 - val_loss:
0.5690 - val_acc: 0.8562 - val_binary_crossentropy: 0.4934
Epoch 20/20
- 5s - loss: 0.1263 - acc: 0.9876 - binary_crossentropy: 0.0509 - val_loss:
0.5766 - val_acc: 0.8548 - val_binary_crossentropy: 0.5014

```

$\lambda^2(0.001)$  means that every weight in the weight matrix will add  $0.001w^2$  to the total loss. Because of this the loss during training is much higher than during testing.

```
[18]: plot_history([('baseline', baseline_history),
                  ('l2', l2_model_history)])
```



### 1.3.2 Dropout

One of the most effective regularization techniques. It consists on dropping out (i.e. set to zero) some output features of the layer during training randomly.

Ex.: A layer output [0.2, 0.5, 1.3, 0.8, 1.1] => [0, 0.5, 1.3, 0, 1.1] (random case)

The “dropout rate” is the fraction of features to be set to zero (usually between [0.2, 0.5]). This is not done during testing, instead outputs are scaled down by the dropout rate to compensate.

```
[13]: dpt_model = keras.models.Sequential([
        keras.layers.Dense(16, activation=tf.nn.relu, input_shape=(NUM_WORDS,)),
        keras.layers.Dropout(0.5),
        keras.layers.Dense(16, activation=tf.nn.relu),
        keras.layers.Dropout(0.5),
        keras.layers.Dense(1, activation=tf.nn.sigmoid)
    ])

dpt_model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy', 'binary_crossentropy'])

dpt_model_history = dpt_model.fit(train_data, train_labels,
                                  epochs=20,
                                  batch_size=512,
                                  validation_data=(test_data, test_labels),
                                  verbose=2)
```

WARNING:tensorflow:From C:\Users\Pedro\AppData\Roaming\Python\Python37\site-packages\tensorflow\python\keras\layers\core.py:143: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Train on 25000 samples, validate on 25000 samples

Epoch 1/20

- 6s - loss: 0.6203 - acc: 0.6558 - binary\_crossentropy: 0.6203 - val\_loss: 0.4907 - val\_acc: 0.8448 - val\_binary\_crossentropy: 0.4907

Epoch 2/20

- 7s - loss: 0.4535 - acc: 0.8092 - binary\_crossentropy: 0.4535 - val\_loss: 0.3468 - val\_acc: 0.8820 - val\_binary\_crossentropy: 0.3468

Epoch 3/20

- 6s - loss: 0.3546 - acc: 0.8693 - binary\_crossentropy: 0.3546 - val\_loss: 0.2948 - val\_acc: 0.8883 - val\_binary\_crossentropy: 0.2948

Epoch 4/20

- 6s - loss: 0.2969 - acc: 0.9010 - binary\_crossentropy: 0.2969 - val\_loss:

0.2782 - val\_acc: 0.8892 - val\_binary\_crossentropy: 0.2782

Epoch 5/20

- 7s - loss: 0.2524 - acc: 0.9196 - binary\_crossentropy: 0.2524 - val\_loss: 0.2805 - val\_acc: 0.8892 - val\_binary\_crossentropy: 0.2805

Epoch 6/20

- 8s - loss: 0.2191 - acc: 0.9317 - binary\_crossentropy: 0.2191 - val\_loss: 0.2879 - val\_acc: 0.8834 - val\_binary\_crossentropy: 0.2879

Epoch 7/20

- 6s - loss: 0.1945 - acc: 0.9401 - binary\_crossentropy: 0.1945 - val\_loss: 0.2971 - val\_acc: 0.8840 - val\_binary\_crossentropy: 0.2971

Epoch 8/20

- 6s - loss: 0.1738 - acc: 0.9464 - binary\_crossentropy: 0.1738 - val\_loss: 0.3103 - val\_acc: 0.8824 - val\_binary\_crossentropy: 0.3103

Epoch 9/20

- 6s - loss: 0.1554 - acc: 0.9532 - binary\_crossentropy: 0.1554 - val\_loss: 0.3375 - val\_acc: 0.8818 - val\_binary\_crossentropy: 0.3375

Epoch 10/20

- 6s - loss: 0.1429 - acc: 0.9562 - binary\_crossentropy: 0.1429 - val\_loss: 0.3451 - val\_acc: 0.8776 - val\_binary\_crossentropy: 0.3451

Epoch 11/20

- 6s - loss: 0.1294 - acc: 0.9615 - binary\_crossentropy: 0.1294 - val\_loss: 0.3725 - val\_acc: 0.8785 - val\_binary\_crossentropy: 0.3725

Epoch 12/20

- 6s - loss: 0.1177 - acc: 0.9644 - binary\_crossentropy: 0.1177 - val\_loss: 0.3877 - val\_acc: 0.8769 - val\_binary\_crossentropy: 0.3877

Epoch 13/20

- 6s - loss: 0.1092 - acc: 0.9668 - binary\_crossentropy: 0.1092 - val\_loss: 0.4198 - val\_acc: 0.8743 - val\_binary\_crossentropy: 0.4198

Epoch 14/20

- 6s - loss: 0.1040 - acc: 0.9680 - binary\_crossentropy: 0.1040 - val\_loss: 0.4274 - val\_acc: 0.8741 - val\_binary\_crossentropy: 0.4274

Epoch 15/20

- 7s - loss: 0.0964 - acc: 0.9712 - binary\_crossentropy: 0.0964 - val\_loss: 0.4645 - val\_acc: 0.8758 - val\_binary\_crossentropy: 0.4645

Epoch 16/20

- 7s - loss: 0.0921 - acc: 0.9723 - binary\_crossentropy: 0.0921 - val\_loss: 0.4812 - val\_acc: 0.8754 - val\_binary\_crossentropy: 0.4812

Epoch 17/20

- 6s - loss: 0.0899 - acc: 0.9740 - binary\_crossentropy: 0.0899 - val\_loss: 0.5100 - val\_acc: 0.8745 - val\_binary\_crossentropy: 0.5100

Epoch 18/20

- 6s - loss: 0.0813 - acc: 0.9767 - binary\_crossentropy: 0.0813 - val\_loss: 0.5266 - val\_acc: 0.8751 - val\_binary\_crossentropy: 0.5266

Epoch 19/20


- 6s - loss: 0.0773 - acc: 0.9782 - binary\_crossentropy: 0.0773 - val\_loss: 0.5228 - val\_acc: 0.8737 - val\_binary\_crossentropy: 0.5228

Epoch 20/20

- 6s - loss: 0.0748 - acc: 0.9785 - binary\_crossentropy: 0.0748 - val\_loss:

0.5404 - val\_acc: 0.8720 - val\_binary\_crossentropy: 0.5404

```
[19]: plot_history([('baseline', baseline_history),  
                  ('dropout', dpt_model_history)])
```



### 1.3.3 Combining both strategies

```
[15]: early_stop = keras.callbacks.EarlyStopping(monitor='val_loss', patience=3)  
  
all_model = keras.models.Sequential([  
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),  
        ↪activation=tf.nn.relu, input_shape=(NUM_WORDS,)),  
    keras.layers.Dropout(0.5),  
    keras.layers.Dense(16, kernel_regularizer=keras.regularizers.l2(0.001),  
        ↪activation=tf.nn.relu),  
    keras.layers.Dropout(0.5),  
    keras.layers.Dense(1, activation=tf.nn.sigmoid)  
)  
  
all_model.compile(optimizer='adam',  
                  loss='binary_crossentropy',  
                  metrics=['accuracy', 'binary_crossentropy'])
```

```
all_model_history = all_model.fit(train_data, train_labels,
                                   epochs=20,
                                   batch_size=512,
                                   validation_data=(test_data, test_labels),
                                   verbose=1)
                                   #callbacks=[early_stop])
```

Train on 25000 samples, validate on 25000 samples

Epoch 1/20

25000/25000 [=====] - 7s 282us/sample - loss: 0.6899 -  
acc: 0.5950 - binary\_crossentropy: 0.6545 - val\_loss: 0.5791 - val\_acc: 0.8408 -  
val\_binary\_crossentropy: 0.5476

Epoch 2/20

25000/25000 [=====] - 7s 286us/sample - loss: 0.5444 -  
acc: 0.7659 - binary\_crossentropy: 0.5110 - val\_loss: 0.4123 - val\_acc: 0.8755 -  
val\_binary\_crossentropy: 0.3762

Epoch 3/20

25000/25000 [=====] - 7s 264us/sample - loss: 0.4426 -  
acc: 0.8405 - binary\_crossentropy: 0.4031 - val\_loss: 0.3540 - val\_acc: 0.8833 -  
val\_binary\_crossentropy: 0.3109

Epoch 4/20

25000/25000 [=====] - 7s 269us/sample - loss: 0.3776 -  
acc: 0.8819 - binary\_crossentropy: 0.3312 - val\_loss: 0.3345 - val\_acc: 0.8878 -  
val\_binary\_crossentropy: 0.2850

Epoch 5/20

25000/25000 [=====] - 6s 258us/sample - loss: 0.3419 -  
acc: 0.9000 - binary\_crossentropy: 0.2895 - val\_loss: 0.3322 - val\_acc: 0.8864 -  
val\_binary\_crossentropy: 0.2770

Epoch 6/20

25000/25000 [=====] - 6s 254us/sample - loss: 0.3133 -  
acc: 0.9144 - binary\_crossentropy: 0.2559 - val\_loss: 0.3395 - val\_acc: 0.8859 -  
val\_binary\_crossentropy: 0.2800

Epoch 7/20

25000/25000 [=====] - 7s 271us/sample - loss: 0.2916 -  
acc: 0.9230 - binary\_crossentropy: 0.2301 - val\_loss: 0.3489 - val\_acc: 0.8840 -  
val\_binary\_crossentropy: 0.2856

Epoch 8/20

25000/25000 [=====] - 7s 261us/sample - loss: 0.2832 -  
acc: 0.9283 - binary\_crossentropy: 0.2183 - val\_loss: 0.3585 - val\_acc: 0.8798 -  
val\_binary\_crossentropy: 0.2920

Epoch 9/20

25000/25000 [=====] - 6s 255us/sample - loss: 0.2699 -  
acc: 0.9321 - binary\_crossentropy: 0.2019 - val\_loss: 0.3733 - val\_acc: 0.8805 -  
val\_binary\_crossentropy: 0.3040

Epoch 10/20

25000/25000 [=====] - 6s 247us/sample - loss: 0.2620 -  
acc: 0.9345 - binary\_crossentropy: 0.1915 - val\_loss: 0.3895 - val\_acc: 0.8796 -



```


val_binary_crossentropy: 0.3178
Epoch 11/20
25000/25000 [=====] - 6s 246us/sample - loss: 0.2519 -
acc: 0.9397 - binary_crossentropy: 0.1790 - val_loss: 0.3855 - val_acc: 0.8785 -
val_binary_crossentropy: 0.3116
Epoch 12/20
25000/25000 [=====] - 6s 257us/sample - loss: 0.2491 -
acc: 0.9415 - binary_crossentropy: 0.1744 - val_loss: 0.3999 - val_acc: 0.8776 -
val_binary_crossentropy: 0.3245
Epoch 13/20
25000/25000 [=====] - 6s 245us/sample - loss: 0.2445 -
acc: 0.9420 - binary_crossentropy: 0.1684 - val_loss: 0.4088 - val_acc: 0.8786 -
val_binary_crossentropy: 0.3321
Epoch 14/20
25000/25000 [=====] - 6s 242us/sample - loss: 0.2412 -
acc: 0.9446 - binary_crossentropy: 0.1637 - val_loss: 0.4132 - val_acc: 0.8784 -
val_binary_crossentropy: 0.3349
Epoch 15/20
25000/25000 [=====] - 6s 254us/sample - loss: 0.2342 -
acc: 0.9472 - binary_crossentropy: 0.1549 - val_loss: 0.4244 - val_acc: 0.8782 -
val_binary_crossentropy: 0.3442
Epoch 16/20
25000/25000 [=====] - 6s 260us/sample - loss: 0.2322 -
acc: 0.9492 - binary_crossentropy: 0.1516 - val_loss: 0.4393 - val_acc: 0.8782 -
val_binary_crossentropy: 0.3583
Epoch 17/20
25000/25000 [=====] - 6s 258us/sample - loss: 0.2268 -
acc: 0.9493 - binary_crossentropy: 0.1454 - val_loss: 0.4342 - val_acc: 0.8745 -
val_binary_crossentropy: 0.3526
Epoch 18/20
25000/25000 [=====] - 7s 263us/sample - loss: 0.2276 -
acc: 0.9475 - binary_crossentropy: 0.1454 - val_loss: 0.4331 - val_acc: 0.8767 -
val_binary_crossentropy: 0.3505
Epoch 19/20
25000/25000 [=====] - 6s 254us/sample - loss: 0.2252 -
acc: 0.9512 - binary_crossentropy: 0.1422 - val_loss: 0.4447 - val_acc: 0.8745 -
val_binary_crossentropy: 0.3613
Epoch 20/20
25000/25000 [=====] - 6s 258us/sample - loss: 0.2239 -
acc: 0.9508 - binary_crossentropy: 0.1402 - val_loss: 0.4495 - val_acc: 0.8749 -
val_binary_crossentropy: 0.3656

```

```

[21]: plot_history([('baseline', baseline_history),
                  ('dropout', dpt_model_history),
                  ('l2', l2_model_history),
                  ('all', all_model_history)])

```



## 1.4 Conclusion

Adding dropout is a clear improvement over the baseline model.

To recap: here the most common ways to prevent overfitting in neural networks:

- Get more training data.
- Reduce the capacity of the network.
- Add weight regularization.
- Add dropout.
- And two important approaches not covered in this notebook are data-augmentation and batch normalization.